

A Supplemental Materials

B Qualitative Timeline Comparisons

This appendix visualizes the final timelines for the *Anesthesia* scene to compare the baseline manual workflow (P1) with *MicroEditor* (P6). As noted in Section 6.2, micro-editing is subjective. Overall pacing and reaction shot placements naturally vary based on individual editing styles. However, when perceptual cues are strong, we observe similar timing outcomes. Figure 9 shows both participants arriving at a comparable J-cut offset (0.9s vs. 1.1s) near the end of the scene, driven by a cut-on-action. To illustrate the underlying method that generates these suggestions, Figure 10 provides a detailed view of the system’s internal synchrony model for the same *Anesthesia* scene. It visualizes the word-level ASR alignments, the continuous audio-visual synchrony curves (σ), and the resulting calculated bounds that constrain the suggestions surfaced in UI.

Similar timeline visualizations and synchrony analyses for other scenes can be found in our interactive supplemental gallery.

C Preprocessing Pipeline

This section details the offline preprocessing that extracts the signals used by the interactive tools described in Section 5. Given source footage *M* and a rough cut *R*, preprocessing runs in three stages before the editing session begins (Figure 4).

C.1 System Implementation and Prompt Details

Since VLMs excel at coarse scene understanding but struggle with micro-temporal reasoning due to their inherent sparse sampling [36], we use VLMs for semantic analysis (character identification, shot classification) while relying on traditional computer vision for frame-level signals (synchrony proxy, blink detection).

Stage 1: VLM Coarse Analysis (2–5 min): A VLM (Gemini) analyzes all source footages sampled at 1 FPS to within 1 API call to do cross-associate semantic information and produce a *global character map*. For each source clip $M_j \in \mathbf{M}$, the VLM identifies distinct characters, shot types, and non-usable regions. We use 480p proxy files to reduce token cost while preserving sufficient detail for character recognition. The prompt structure:

Analyze ALL footage files together to identify characters and shot types.

CHARACTER IDENTIFICATION: *Identify all distinct people and objects of interest across all footage files (Figure 11). The same person in different footage = same character_id. Look for: hair, clothing, glasses, facial features, body type. A close-up in one footage and wide shot in another can be the same person.*

SHOT TYPE CLASSIFICATION: *Classify each segment as: extreme_close_up, close_up, medium_close_up, medium, medium_wide, wide, over_the_shoulder, two_shot, or group_shot.*

OVER-THE-SHOULDER RULE: *In OTS shots, assign character_id to the person whose face is visible (not the person shown from behind). This ensures visible faces for reaction shot extraction.*

NON-USABLE SEGMENTS: *Mark slates, crew visible, camera movement, and empty frames as no_character.*

TEMPORAL COVERAGE: *Cover the entire duration of each footage file with no gaps. Approximate timing (± 1 –2 seconds) is acceptable.*




Figure 9: Two participants editing the same scene *Anesthesia*: P6 with MicroEdits (top) and P1 editing manually (bottom). The MicroEdits row labels cut types per user’s decisions (e.g., optimization based on AV synchrony, cut-on-action). When cut-on-action is a strong cue, manual and assisted edits sometimes converge on similar cut timing (0.9 vs 1.1s J-cut offset).

Stage 2: AV Sync Processing (2–5 min): In parallel, we extract frame-level signals locally:

- **Word boundaries:** ASR (Parakeet-TDT [41] or Whisper [37]) provides word-level timestamps, defining bounds for rhythm editing windows.
- **Synchrony function:** For dialogue scenes, active speaker detection (LR-ASD [21]) computes per-frame lip-audio correlation for each face track. For non-dialogue content (e.g., sports, music), we use an audio-visual alignment network (AVID-CMA [25]). We threshold the continuous signal to obtain $\sigma(M_j, t) \in \{0, 1\}$.
- **Blink onsets:** MediaPipe [22] estimates eye aspect ratio per frame; blink onset occurs when EAR falls below threshold.

Stage 3: VLM-ASD Correspondence (instant): We link VLM character identities to ASD face tracks by extracting representative face crops from each track and querying Gemini to match them against VLM-generated character thumbnails. When matching is ambiguous, we fall back to DeepFace embeddings (Facenet512) with cosine similarity. This enables per-character queries of σ and supports the character focus optimization (Section 5.3).

C.2 Cut-on-Action Detection

Current state of VLMs cannot reliably reason about frame-level motion continuity [36]. We address this by converting the temporal reasoning problem into a spatial matching task. The model is prompted with the 2x2 grid layout (sample from the left and right bounds of *W*) and temporal alignment rules:

You are a professional film editor with expertise in cut-on-action techniques.

GRID LAYOUT: *Top row shows earlier cut opportunities (outgoing end, incoming preroll). Bottom row shows later cut opportunities (outgoing post-roll, incoming start).*




Figure 10: System-level synchrony analysis for an edit of *Anesthesia* scene. The visualization shows the continuous synchrony signal (σ , purple line), word-level timing (yellow boxes), and the calculated safe decoupling windows used to generate the L/J-cut and Cut-on-Action suggestions surfaced to the user.

TEMPORAL ALIGNMENT RULES:

- For later cuts (J-cut): Match frames with equal numbers ($A=B$), since frame 1 = cut point for both grids.
- For earlier cuts (L-cut): Match frames by distance from last frame, maintaining consistent offset ($B-A$).

FOCUS ON: Motion trajectory, direction continuity, body alignment (head angle, shoulder rotation, limb positions), and temporal alignment.

The temporal alignment rules encode the key insight that frames are sampled at the same rate (4.5 fps), so valid cut points must match frames representing the same moment relative to the cut boundary.

Motion Matching: The VLM identifies frame pairs where motion trajectory, body position, or gaze continues smoothly across the cut boundary. Each pair (n_{out}, n_{in}) converts to a timeline offset:

$$\delta = (n_{frame} - 1) / \text{fps}_{\text{mosaic}}$$




Figure 11: View identification using vision-language model analysis across different content types. Left: Dialogue scene with two distinct character views identified by visual features (hair, clothing patterns). Right: Product review distinguishing human presenter views from product demonstration shots. The VLM extracts distinctive features for each view, enabling the View Focus feature to redistribute screen time through automated L/J cut suggestions.

Top-row matches yield L-cuts ($\delta < 0$). Bottom-row matches yield J-cuts ($\delta > 0$).

C.3 Reaction Shot Indexing

Low-sync segments provide source material for reaction shots. A segment from M_j qualifies as a candidate if:

- $\sigma(M_j, t) = 0$ throughout (character visible but not speaking)
- The target character is detected in $\geq 80\%$ of frames
- The segment is adjacent to spoken material (filtering irrelevant silent stretches)

We index candidates by character identity and emotion embeddings (EmoCLIP [15]), enabling the semantic search interface described in Section 5.4.

C.4 Processing Performance

Table 3 shows typical times for a one-minute dialogue scene with 4 source clips (~5 minutes raw footage).

Stages 1 and 2 run in parallel; VLM API calls are the primary bottleneck. All signals are computed in source time and remain stable across timeline changes, so preprocessing runs once per project.

D VLM Input Strategies for Match-on-Action Detection

We explore different input strategies for using vision-language models (VLMs) to detect match-on-action cut points. Our goal is to identify an efficient representation that balances accuracy with inference speed for interactive editing workflows.

Table 3: Preprocessing time breakdown (1-minute scene, 4 source clips, MacBook Pro M2)

Stage / Component	Time
<i>Stage 1: VLM Coarse Analysis</i>	
Proxy generation (480p)	30s
VLM footage analysis (Gemini)	90s
<i>Stage 2: AV Sync Processing</i>	
ASR transcription	45s
Active speaker detection (LR-ASD)	60s
Blink detection (MediaPipe)	15s
<i>Stage 3: VLM-ASD Matching</i>	
Face embedding + matching	20s
<i>Per-Cut Processing</i>	
Cut-on-action mosaics + VLM	60s
Total	~5 min

D.1 Input Representations

For each potential cut point, 24-frame sequences are extracted from the outgoing shot’s post-roll (unused footage beyond the cut) and incoming shot’s pre-roll (footage before timeline start) at 4.5 fps. Frames are arranged into 4×6 grid mosaics (Figure 14).

We compare two input layouts: (1) **2×2 Grid**, a single mosaic with four sequences generating 6 suggestions in one API call; (2) **Two-Image**, separate mosaics for outgoing and incoming sequences generating 3 suggestions in one API call. To analyze both cut directions, the two-image strategy requires two API calls ($2 \times$ inference time).






Figure 12: Micro-edits & Their Impact on Rhythm: A cut from source media M_1 to source media M_2 can be realized in multiple ways, starting from $M_1(S_1)$. The figure contrasts four strategies—initial edit, time-preserving edit, time-contracting edit, and time-expanding edit—each altering the pacing and perceived rhythm of the sequence by adjusting how much time is borrowed or extended around the cut.


TOP ROW: Earlier Cuts | Match-on-Action A→B (offset = B−A)
A: Outgoing (end of shot)



B: Incoming (preroll)



BOTTOM ROW: Later Cuts | Match-on-Action A→B (A = B)
A: Outgoing (postroll)



B: Incoming (start of shot)




Figure 13: Cut-on-action with VLM explanations for an L-cut (top) and J-cut (bottom). L-cut: “LEFT: B21→A2. Stepping back one frame... The man’s slight head tilt is consistent between the close-up and the wide shot.” J-cut: “RIGHT: A4→B4. This is the moment the camera whips/pans... The blur in the outgoing shot matches the motion blur and head turn initiation in the incoming shot, hiding the cut effectively.”

D.2 Comparison

Ground truth is established by a professional editor. For J-cuts, valid matches require temporal alignment ($A=B$) with matching frames in the 4-6 range where motion continuity is strongest. For L-cuts, valid matches require consistent temporal offset ($B-A$) across suggestions. We measure success rate (percentage of runs identifying at least one ground truth match) and inference time. We compared three input strategies with Gemini 3.0 Flash (Table 4).

All strategies achieved 100% temporal alignment ($A=B$ for J-cuts, consistent $B-A$ offset for L-cuts). The combined 2×2 grid provides the fastest inference (3.84s) while maintaining equivalent accuracy, making it well-suited for interactive editing workflows. We use this strategy in our system.

Table 4: Gemini 3.0 Flash performance across input strategies.

Strategy	Success	Time (s)	Tokens
Combined (2×2 grid)	100% (10/10)	3.84	2,389
Split (two images)	100% (10/10)	5.76	5,647
Video	100% (10/10)	8.36	5,878

E Full Formative Study Quotes and Details

To better understand the micro-editing process, we recruited six professional video editors (5-15 years of experience). Participants edited two short dialogue scenes (approximately 1 minute each): a conversation between a nursing home manager and a resident



(a) 2×2 grid strategy: Single mosaic containing four sequences (earlier and later cut opportunities for both outgoing and incoming footage). Model analyzes entire grid in one API call, generating 6 frame-pair suggestions.



(b) Two-image strategy: Separate mosaics for outgoing and incoming sequences sent in one API call, generating 3 frame-pair suggestions.

Figure 14: Visual input representations for match-on-action detection. Each mosaic shows 24 sequential frames numbered 1-24 from left to right, top to bottom. Test sequence depicts a body turn requiring precise temporal alignment for smooth continuity.

being kicked out (“Bingo Night”), and a lighter interaction in an antique shop (“Soup Spoon”). Each participant completed four editing tasks per scene with different objectives: (1) refining timing using only straight cuts, (2) converting specific cuts to J-cuts only, (3) converting to L-cuts only, and (4) creating their “best edit” using any combination of techniques. Clip selection is not part of the task.

Users are only adjusting cuts. Participants used their preferred editing software (3 Adobe Premiere Pro, 2 DaVinci Resolve, 1 Final Cut Pro) while thinking aloud about their decision-making process. We found that editors often conceptualize editing along two primary dimensions: **rhythm** (temporal flow) and **view** (visual continuity).

E.1 Observation 1: Rhythm Edits for Isolating Temporal Flow

Editors describe audio and video as “two different systems” (P1) that can be manipulated independently to create different psychological effects. As P1 articulated: “What you’re interested in looking at is not the same as what you’re interested in hearing. So if you introduce the hearing as one sense and then introduce the sight, it’s a different feeling than if you introduce both at the same time.”

E.1.1 Temporal Focus Through Isolation. Half of our participants began with a *radio edit*: cutting dialogue for natural conversation flow before addressing visuals. P4 described this isolation strategy: “I’m going to close my eyes and listen...I was trying to listen to it as if it sounded like a natural conversation....this lets me really focus on the flow and not be interrupted by some video breaks.” Similarly, P5 noted that in dialogue scenes, “Even if the visuals look nice and flow together well, if the audio feels off, the whole scene feels wrong.” This “listening” is less about audio analysis and more about feeling time and the pacing and rhythm of conversation. As P1 noted, awkward pacing is easily noticeable: “Beginner editors often... leave too long of a pause after a line or not enough of a pause.” During this stage, all participants actively shaped the audio timeline, shifted straight cuts to remove dead air, trimmed silences, and ripple-deleted gaps. They repeatedly replayed transition, nudging cut points by a few frames until the pacing felt right.

E.1.2 Preserving Established Timing. Once editors established temporal flow, they actively avoided disrupting it during later visual refinements. They manipulated the video tracks by trimming and ripple-adjusting visuals, adding L- or J-cuts, or extending reaction shots, without making major modifications to the underlying dialogue rhythm. P1 emphasized “I already changed the timing... I don’t want to mess with that.” This protective behavior reflects the cognitive challenge of managing what Walter Murch calls the six criteria of editing (emotion, story, rhythm, eye-trace, planarity, and spatial continuity) simultaneously [26]. When editors adjust both rhythm and visuals together, each change can cascade into unintended consequences: fixing a visual match might disrupt carefully crafted pacing, while adjusting timing might break visual continuity. While Murch suggests editors should prioritize emotion over technical continuity when forced to choose, our participants revealed a complementary strategy: separating temporal decisions (rhythm, emotion, story) from visual ones (eye-trace, continuity) into distinct passes. This pattern echoes established editing texts [26, 35]: rhythm decisions precede and provide structure for visual refinement. While these timing decisions are intuitive, what P3 called “vibe editing”, they form the structure that subsequent visual editing should preserve.

E.2 Observation 2: View Edits and Navigating an Overwhelming Decision Space

After establishing rhythm, editors shift focus to visual refinement, a phase characterized by a tension between frame-level precision and decision paralysis.

E.2.1 Frame-level scrubbing. Editors need frame-level control, but sometimes induce cognitive load of many possibilities, as P4 articulated “I just have this bias of like, oh, I made an edit. Now I don’t want to touch it too much. Otherwise, too many capabilities are getting open.” Even when editors embrace frame-level exploration, the process lacks structure. P5 describes the typical workflow: “I’ll typically fiddle around with it like going frame by frame, removing one frame here, removing a few frames here, adding back a few frames here, just to see what looks best visually.” This trial-and-error approach is time-consuming and often yields diminishing returns, with P4 noting that “there’s a lot of like, editors moving one frame like perfect... it’s just in our brain.”

E.2.2 Structuring decisions through visual goals. Despite this complexity, editors consistently evaluate cuts against specific visual criteria. Our participants identified several recurring considerations that guide their frame-level decisions: guiding viewer attention, preserving action continuity, and avoiding cuts on blinks.

L/J cuts for Guiding Viewer Attention: Editors use split edits (L/J cuts) to guide viewer attention, and P1 offered a compelling perceptual explanation: “If you’re from the perspective of someone who’s listening to people talk... You’re thinking about this person talking. Then another person talks, and they might get in some words before you turn around and look at them. I think that’s what a J-cut is.” J-cuts were often used to anticipate dialogue or smooth transitions, as P1 explained: “J is for reaction to the talking... They start talking. And you...haven’t quite moved on from looking at this person.” Conversely, L-cuts allow editors to extend reactions and “let a moment breathe,” with P4 noting, “For L-cuts, it’s mostly about seeing other people’s reactions. I like this extended reaction where he speaks... I like seeing her face a little bit more.” When both P1 and P2 repeatedly express wanting to “see her [the senior lady in Bingo Night] more,” they agree that the person speaking isn’t always the person we need to see. Split edits enable this redistribution of visual attention, and replicate natural attention patterns in real conversation, where our visual focus lags behind audio cues.

Cut on Action: All participants sought moments where movement could mark transitions. P1 emphasized, “There also needs to be visual continuity... if she... is shaking her head in this... and then at the beginning of this shot she’s no longer shaking her head...She’s looking at a different direction. That doesn’t make any sense narratively; it’ll throw the viewer off.” Similarly, P5 described: “When he starts moving in this cut, it carries over to the next cut... the action is the main thing I was editing based on.” These action matches are discrete decision points within the continuous timeline, yet identifying them consumed significant portions of our participants’ 1.5-2 hours editing sessions.

Cut before Blinks: Several participants considered blinks when choosing cut points, referencing Walter Murch’s *In the Blink of an Eye*. P3 noted “I was always told you shouldn’t end a cut on a blink,” while P5 explained, “If someone blinks and you cut on that blink, you could just tell...it immediately draws you to the fact that this is an edit.” Moments just before a blink represent good cut opportunities, but they are easy to miss during manual review.

These findings present both a problem and an opportunity. While editors have clear visual goals (character focus, action matching, cut before blinks), manually identifying these moments across even short scenes required 1.5-2 hours of frame-by-frame exploration. Though editors routinely use markers to annotate important moments, finding these moments remains manual. This suggests potential for computational assistance in automatically identifying these decision points, structuring the overwhelming continuous timeline into discrete, meaningful choices.

F User Study 1 Details

Building on the insights identified in our formative study (Section 3), we evaluated how MicroEdits supports (1) **rhythm** (temporal flow) and (2) **view** (visual continuity and narrative focus) while reducing cognitive load and encouraging creative iteration.

F.1 Participants

We recruited 6 professional video editors (3 women, 3 men, aged 18-27) with 5-15 years of editing experience (Table 5). One edits daily (P3), three edit weekly or multiple times per week (P1, P6 edit weekly; P3 daily), two edit monthly (P2, P5), and one edits yearly (P4). Participants work across narrative films (5 editors), documentaries (3 editors), YouTube content (2 editors), and social media content (2 editors). All primarily use Adobe Premiere Pro in their professional work.

F.2 Study Format

All participants completed a 60-minute in-person session. After a tutorial on MicroEdits terminology and features, participants edited two 1-minute dialogue scenes using our interface: “Pumpkin Princess” (awkward romantic interaction between a nerdy outcast and her crush) and “Anesthesia” (suspenseful 1950s scene where a housewife suspects her dentist). Participants completed both conditions in counterbalanced order: manual editing (all features disabled) and assisted editing (features enabled). We encouraged think-aloud protocol throughout. Following each edit, participants rated features on 7-point Likert scales and completed creativity support metrics adapted from the Creativity Support Index 6.

G Full User Study 2 Details

Building on User Study 1 (Sec F), we conducted a second evaluation to assess the newly added **Reaction Shot Selection** feature with emotion filtering, while continuing to evaluate the core rhythm and view pass capabilities with a larger participant pool.

G.1 Participants

We recruited 10 video editors (6 women, 4 men, aged 19-28) with 1-10 years of editing experience (Table 7). Participants edit at varying frequencies: two edit weekly (P2, P8), three edit monthly (P4, P7, P9), four edit every few weeks to monthly (P1, P3, P5, P6), and one edits a few times per year (P10). Participants work across narrative films (7 editors), commercial/social media content (5 editors), documentaries (3 editors), and YouTube/digital content (3 editors). Most primarily use Adobe Premiere Pro, with some using DaVinci Resolve or Final Cut Pro.

G.2 Study Format

All participants completed a 60-minute in-person session following the same protocol as User Study 1. After a tutorial covering MicroEdits features including the new Reaction Shot Selection tool, participants edited two 1-minute dialogue scenes (“Pumpkin Princess” and “Anesthesia”) in counterbalanced order across manual and assisted conditions. Following each edit, participants rated features on 7-point Likert scales and completed creativity support metrics (Table 8). Additionally, participants rated the degree to which pacing and visual flow improved over the original rough assembly.

H Synchrony Study Details

For each of the 14 video clips, participants on Prolific.com were shown the embedded video player and asked the following two questions. The questions were repeated identically for every clip.

Q1: Did you notice any distracting audio-visual mismatches or lip-sync errors in this clip?

- Yes
- No

Q2: Overall, how natural did the flow of this clip feel?

- 1 - Very jarring / distracting
- 2
- 3
- 4 - Neutral
- 5
- 6
- 7 - Very smooth / natural

Table 5: Participant demographics and feature ratings from the MicroEdits user study

ID	Demographics				MicroEdits Feature Ratings			
	Age	Gender	Experience	Content Types	Rhythm	Character Focus	Action	Blink
P1	22	F	7+ years, Premiere	Narrative, YouTube, TikTok	6	7	5	7
P2	27	M	15 years	Narrative, Documentary	6	6	7	5
P3	26	F	7 years, Premiere	Narrative, Documentary	7	7	6	4
P4	26	F	5 years	Narrative, YouTube	6	7	7	4
P5	18	M	6 years	Narrative, Digital Media	6	7	7	6
P6	19	M	8 years	Narrative, Documentary	6	6	7	7

Note: Participants completed two editing tasks using the MicroEdits interface in counterbalanced order: manual editing (all features disabled) and assisted editing (features enabled). Scene assignment (Pumpkin Princess/Anesthesia) and condition order were counterbalanced. Feature ratings on 1-7 scale where 7 = Extremely Useful. Rhythm = Pre-roll & Trim, Character Focus = Emphasize Speaker (automatic L/J cuts), Action = Cut-on-action suggestions, Blink = Blink detection markers.

Table 6: Creativity support and user experience metrics. Participants rated agreement on a 7-point scale (1=Strongly Disagree, 7=Strongly Agree) with statements adapted from the Creativity Support Index and task-specific measures.

Metric	P1	P2	P3	P4	P5	P6	Median
<i>Creativity Support Dimensions</i>							
Exploration: Helped explore possibilities	6	7	6	6	6	5	6.0
Exploration: Encouraged different cuts	7	7	7	6	6	4	6.5
Expressiveness: Achieved creative vision	5	5	6	6	6	6	6.0
Expressiveness: Achieved editing style	6	5	6	6	5	5	5.5
Effort/Reward: Quality worth effort	6	6	6	6	6	4	6.0
Results Worth Effort: Satisfied with output	5	5	6	6	6	7	6.0
Immersion: Attention on task not tool	5	5	6	5	7	6	5.5
Enjoyment: Enjoyed using system	6	7	7	7	7	6	7.0
<i>Task-Specific Efficiency</i>							
Decision Support: Better editing decisions	6	7	7	6	6	5	6.0
Confidence: More confident in choices	6	5	5	6	5	4	5.0
Efficiency: Saved time	5	7	7	7	7	6	7.0

Table 7: Participant demographics and feature ratings from User Study 2 (with Reaction Shot Retrieval)

ID	Demographics				MicroEdits Feature Ratings				
	Age	Gender	Experience	Content Types	Rhythm	Char. Focus	Action	Blink	Reaction
P1	20	F	10+ years	Narrative, Animation	6	6	7	4	6
P2	20	F	3 years	Dance, Trailers	6	7	6	5	7
P3	20	M	3-4 years	Reels, Product Videos	5	6	4	1	7
P4	24	M	7 years	Narrative, News	7	6	7	5	7
P5	19	F	4 years	Narrative, Commercial	6	7	6	3	7
P6	27	F	3+ years	Narrative, Architectural	7	6	7	5	7
P7	27	M	10 years	Narrative	6	5	6	5	7
P8	24	F	1 year	Short Reels/TikTok	7	7	7	4	7
P9	25	F	3 years	Music Videos, Interviews	5	6	5	5	7
P10	28	M	10 years	Narrative, Documentary	6	6	7	4	7

Note: Feature ratings on 1-7 scale where 7 = Extremely Useful. Rhythm = Pacing Control via Trim/Extend, Char. Focus = Character Focus with automated L/J cuts, Action = Cut-on-action suggestions, Blink = Blink detection markers, Reaction = Reaction Shot Selection with Emotion Filtering.

Table 8: Creativity support and user experience metrics for User Study 2. Participants rated agreement on a 7-point scale (1=Strongly Disagree, 7=Strongly Agree).

Metric	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10	Med.
<i>Creativity Support Dimensions</i>											
Exploration: Helped explore possibilities	6	6	5	7	7	7	6	6	6	6	6.0
Exploration: Encouraged different cuts	7	5	6	6	7	7	6	6	6	6	6.0
Expressiveness: Achieved creative vision	5	5	6	4	5	7	5	6	5	6	5.0
Expressiveness: Achieved editing style	5	6	6	5	6	7	6	6	5	6	6.0
Effort/Reward: Quality worth effort	6	6	5	4	7	7	6	6	5	7	6.0
Results Worth Effort: Satisfied with output	6	7	6	6	7	7	6	6	4	7	6.0
Immersion: Attention on task not tool	5	6	4	7	5	6	7	6	3	5	5.5
Enjoyment: Enjoyed using system	6	7	5	7	7	7	6	7	5	7	7.0
<i>Task-Specific Efficiency</i>											
Decision Support: Better editing decisions	6	6	6	6	7	7	6	7	4	6	6.0
Confidence: More confident in choices	5	7	6	5	7	7	5	6	5	6	6.0
Efficiency: Saved time	7	7	6	7	7	7	7	7	5	7	7.0
<i>Additional Rating</i>											
Pacing/Flow: Improved over rough assembly	7	6	6	6	6	7	6	6	5	7	6.0